


Structural recovery vs numeric vocabulary size K ($K=24$: <NUM> placeholder; $K=69$: 1-digit; $K=335$: 2-digit; $K=2,418$: 4-digit / current model)

(a) Per-token structural recovery



(b) Per-window structural recovery

