

Per-(axis, slot) recovery vs. tokenized entropy ($n = 30$, patched)

